

# SQOOP COMMANDS

Sqoop:

-->It is a tool which is used to transfer data from relational database to hadoop(hdfs, hive, hbase) and vice versa.

It is used for data ingestion/migration.

Architecture: It is the system design which is used to build a product/software.

Framework: It is used on top of architecture to provide functionality and build the product. It is a combination of tools and components which work together.

Tool: It is an instrument which is used to handle a particular functionality in a software

-->Sqoop IMPORT is used to transfer data from database to hadoop

Sqoop EXPORT is used to transfer data from hadoop to database

Sqoop EVAL is used to get a feel of data ie run queries on tables of databases

-->**To import all the tables**

Sqoop-import-all-tables \

--connect "jdbc:mysql://quickstart.cloudera:3306/retail\_db" \

--username retail\_dba \

--password cloudera \

--as-sequentialfile \ (we mention the file format here default is text file)

--warehouse-dir \output (this directory shdn't exist before)

**Target-dir vs warehouse-dir:**

-->target-dir : It is used when we have a single table to be imported which gets created in one particular folder

--target-dir /user/cloudera/data

--warehouse-dir /user/cloudera/data/orders (if warehouse used for single table)

-->warehouse-dir: it is used to create sub-directories of multiple tables that are imported

--warehouse-dir /user/cloudera/sqoop/employee

--warehouse-dir /user/cloudera/sqoop/persons

--warehouse-dir /user/cloudera/sqoop/retail

--warehouse-dir /user/cloudera/sqoop/billing

## **Sqoop Guide:**

-->sqoop help : Lists all related commands being used in sqoop

-->sqoop help eval : Lists all related commands in "eval"

Redirecting Logs:

### **How to view the job logs later?**

-->If suppose we want to see the logs of the import later we can redirect them to a path where they will be saved and can be viewed later.

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table customers \  
--warehouse-dir /user/cloudera/output1 1>output.log(execution messages)  
2>others.log (complete processing log)
```

**\*\*Note:** The log files can be with any name and extension and stored in local path(/home/cloudera)

-->to view the logs we have to give

-->cat output.log ()

-->cat others.log (where we have complete processing)

-->These log files are stored in local

## **Boundary query:**

### **How mappers divide the work?**

-->To know how 4 mappers divide the work among themselves.

-->It is based on the range defined by the primary key.

[ (max(primary\_key) - min(primary\_key) / 4 ] ~ no. of records taken by each mapper.

-->you can view the runtime log generated where the splitting is taking place

-->based on the range obtained mappers will process data within the range

-->ex. 1-100 : records in range of 1-100 (primary key) is taken up by mapper 1

100-200 : mapper 2

200-300: mapper 3  
300-400: mapper 4

### **How our data can be compressed?**

How we can compress the data ?

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table customers \  
--compress \  
--warehouse-dir /user/cloudera/output1
```

### **How to import only particular columns?**

```
-->To view top 10 records  
-->hadoop fs -ls /output/partm-00000 | head  
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table customers \  
--columns customer_id,customer_fname \  
--warehouse-dir /user/cloudera/output1
```

Putting it all together

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table customers \  
--compress \  
--columns customer_id,customer_fname \  
--where "customer_city in ('new york') " \  
--warehouse-dir /user/cloudera/output1
```

```
--warehouse-dir /user/cloudera/output1
```

### **How to set customized boundary value when we have outliers?**

-->suppose we inserted a record with primary key 800000 where total records is 65000(max primary key)

So system creates boundary value of (800000 - 0 ) which is wrong

Due to this major effort falls on mapper 1 as it processes ~2lakh records in our case all the 65000 records

To avoid this we set the boundary query value

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table customers \  
--boundary-query "select 1, 65000" (here we r specifying the range of primary key)  
--warehouse-dir /user/cloudera/output1
```

### **Split-by USE CASE:**

#### **What if primary key is not present or unevenly distributed?**

-->We can use split-by when a table doesn't have primary keys

-->And when the primary is unevenly distributed

-->we can set mappers to 1 (not recommended as parallelism doesn't exist)

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table customers \  
--split-by "customer_zipcode" \  
--warehouse-dir /user/cloudera/no_pk
```

-->here we have used customer zipcode as index to split data into 4 mappers

-->It creates splits based on max , min values  
-->if the primary key is a text column (customer name) based on ASCII values data will be split.

### **When primary key is text field?**

sqoop import \

**-Dorg.apache.sqoop.splitter.allow\_text\_splitter=true \** (optional as it works without this syntax as well)

--connect "jdbc:mysql://quickstart.cloudera:3306/retail\_db" \

--username retail\_dba \

--password cloudera \

--table customers \

--split-by "customer\_state" \

--target-dir /user/cloudera/no\_pk1

### **How to auto-reset mappers?**

-->It means if there is no primary key then auto set mappers to 1

-->if there is primary key then set mappers to 8

-->by default mappers will be 4

sqoop import \

--connect "jdbc:mysql://quickstart.cloudera:3306/retail\_db" \

--username retail\_dba \

--password cloudera \

--table customers \

**--autoreset-to-one-mapper (or) --m 1 \**

**--num-mappers 8 \**

--target-dir /user/cloudera/no\_pk1

### **How to delimit fields and rows in sqoop?**

-->It means we can separate fields by a delimiter character and separate rows by a delimiter character

sqoop import \

--connect "jdbc:mysql://quickstart.cloudera:3306/retail\_db" \

--username retail\_dba \

```
--password cloudera \  
--table customers \  
--fields-terminated-by '|' \  
--lines-terminated-by ';' \  
--target-dir /user/cloudera/no_pk1
```

### **How to create a hive table using sqoop?**

-->It creates an empty table in hive using sqoop and imports the schema of table from mysql

#### **sqoop create-hive-table \**

```
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table customers \  
--hive-table customers
```

### **How to view detailed logs in sqoop?**

-->"verbose" command is used to view detailed logs of data imported through sqoop which can be used for debugging

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table customers \  
--verbose  
--target-dir /user/cloudera/no_pk1
```

### **How to store output files in already existing folder?**

-->"append" command is used to add output files to a pre-existing folder

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table customers \  
--append
```

```
--target-dir /user/cloudera/no_pk1  
--append
```

### **How to overwrite an output folder?**

-->It causes the existing folder to be deleted (if any) and creates a new folder in its place.

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table customers \  
--target-dir /user/cloudera/no_pk1  
--delete-target-dir
```

### **How to deal with null values?**

-->we can automatically append null values with a predefined value

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_db  
--password cloudera \  
--table customers \  
--target-dir /user/cloudera/no_pk1  
--null-non-string "-1"  
--null-as-string "NA"
```

### **Sqoop Export:**

Sqoop Export:

#### **How to export data from Hdfs to Rdbms using sqoop?**

Step 1: create a table having same schema in rdbms

Step 2: Use sqoop cmd to export the data

```
sqoop export \  
--connect "jdbc:mysql://quickstart.cloudera:3306/banking" \  

```



```
--username root \  
--password cloudera \  
--table card_transactions \  
--export-dir /user/cloudera/card_trans.csv
```

-->data from csv in hadoop is exported to table in msql

### **What happens when we have partial export?**

-->in some use cases when the export fails, then we see that there is partial export of data which is bad for our application.

-->we make sure the database transaction is atomic ie it either succeeds or fails completely.

-->we make use of **staging table** for this purpose which acts like an intermediary.

-->when the export starts records are first fetched into staging table, when there is any error data will not be migrated to main table.

-->if the export succeeds data will be migrated from staging table to main table

```
sqoop export \  
--connect "jdbc:mysql://quickstart.cloudera:3306/banking" \  
--username root \  
--password cloudera \  
--table card_transactions \  
--staging-table card_transactions_stage \  
--export-dir /user/cloudera/card_trans.csv
```

### **Sqoop Incremental:**

#### **What to do when we want to insert new data into hdfs on daily basis?**

There are two methods to do it:

- 1.incremental append
- 2.incremental lastmodified

-->we can insert new data records into hdfs through sqoop regularly.

-->It can be done using the sqoop command

--incremental-append : It is used to insert the new data records into hdfs using sqoop

--check-column : based on which column we can decide on new records to be added (generally primary key)  
--last-value : signifies after this value(primary key) new records can be inserted

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table orders \  
--warehouse-dir /user/cloudera/incremental \  
--incremental append \  
--check-column order_id \  
--last-value 0
```

-->After first import is done and 68883 records were added. Now for the next newly added records

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table orders \  
--incremental-append \  
--check-column order_id \  
--target-dir /user/cloudera/no_pk1 \  
--last-value 68883
```

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table orders \  
--where "order_id > 68000" \  
--warehouse-dir /user/cloudera/incremental \  
--incremental append \  
--check-column order_customer_id \  
--last-value 8440
```

### **How to update and insert records using sqoop?**

-->"incremental lastmodified" is a method used to update the records and insert new records as well.

-->Based on the last modified date we will perform update and insert

-->check for order\_date and update/insert records last modified after '2021-06-25 12:00:00'

-->"append" is used to write the files into given directory without creating new one

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table orders \  
--warehouse-dir /user/cloudera/incremental \  
--incremental lastmodified \  
--check-column order_date \  
--last-value '2021-06-25 12:00:00' \  
--append
```

-->we have to make note of last value so that in our next import we can use it to make new updations/insertion

### **How we can have only latest records in our table?**

-->we get duplicates in our table with both old and new records after updating them.

-->To make sure we have only the latest records in our table we use "merge-key" field in sqoop

-->based on "order\_id" records will be merged and we have only the latest record in our table

```
sqoop import \  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username retail_dba \  
--password cloudera \  
--table orders \  
--warehouse-dir /user/cloudera/incremental \  
--incremental lastmodified \  
--check-column order_date \  
--last-value 0 \ (right from the beginning)  
--merge-key order_id
```

### **Sqoop job:**

#### **How we can automate sqoop import/export?**

-->we can automate running of sqoop jobs using airflow or oozie tools.  
-->As seen previously we have to keep track of last value which has to be passed to the next import which is a tiring task.  
-->we can automate this by creating a sqoop job which will save the state of job and pass it to the next sqoop import of same job.

Job creation:

```
sqoop job \  
--create job_orders \  
-- import \ (there is a space between "--" and "import"  
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \  
--username root \  
--password cloudera \  
--table orders \  
--warehouse-dir /user/cloudera/sqoop_job \  
--incremental append \  
--check-column order_id \  
--last-value 0
```

View list of jobs:

-->sqoop job --list

Run Job:

-->sqoop job --exec job\_orders

To view metadata of job:

-->sqoop job --show job\_orders

### **Where is this metadata stored?**

-->it is stored in /home/cloudera (local path) in a hidden directory (.sqoop)

-->ls -altr

### **\*\*Note:**

-->If the split-by column is a varchar then based on ASCII values the data is divided among the mappers.

-->If the primary key contains randomized values (not exist in real scenario) then u can insert a column with sequential key and make it primary key.

-->If data contains duplicate records even then data is divided based on boundary values query and few mappers may get more data which might be duplicates.

### **-->How boundary value works for outliers**

Suppose ur records are 1,100,200 only three records

Now  $200 - 1/4$  approximately 50 records

1 mapper will only the records between 1 to 50 ie only one record

2nd mapper also holds one

3rd mapper is empty

4th one one record

**Records goes to mapper based on range of primary key**